

```
#IMPORTING LIBRARIES
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math as math
import seaborn as sns
```

```
#LOADING DATA FROM CSV FILE
```

```
import requests
import pandas as pd
import io
```

```
# Corrected URL to point to the raw CSV file
```

```
url = "https://raw.githubusercontent.com/satyanarayanan102/OIBSIP/main/EDA%20Retail%20Sales/"
response = requests.get(url)
```

```
# Check if the request was successful
```

```
if response.status_code == 200:
```

```
    data = pd.read_csv(io.StringIO(response.text))
```

```
    print(data.head()) # Print first few rows to verify
```

```
else:
```

```
    print("Failed to download the file. Check the URL.")
```



	Transaction ID	Date	Customer ID	Gender	Age	Product Category \
0	1	24-11-2023	CUST001	Male	34	Beauty
1	2	27-02-2023	CUST002	Female	26	Clothing
2	3	13-01-2023	CUST003	Male	50	Electronics
3	4	21-05-2023	CUST004	Male	37	Clothing
4	5	06-05-2023	CUST005	Male	30	Beauty

	Quantity	Price per Unit	Total Amount
0	3	50	150
1	2	500	1000
2	1	30	30
3	1	500	500
4	2	50	100

```
#showing the rows
```

```
data.head(1001)
```



	Transaction ID	Date	Customer ID	Gender	Age	Product Category	Quantity	Price per Unit	Total Amount
0	1	24-11-2023	CUST001	Male	34	Beauty	3	50	150
1	2	27-02-2023	CUST002	Female	26	Clothing	2	500	1000
2	3	13-01-2023	CUST003	Male	50	Electronics	1	30	30
3	4	21-05-2023	CUST004	Male	37	Clothing	1	500	500
4	5	06-05-2023	CUST005	Male	30	Beauty	2	50	100
...

```
#gives (num rows, num col)
data.shape
```



```
(1000, 9)
```

```
data.dtypes
```



```
Transaction ID      int64
Date                object
Customer ID         object
Gender              object
Age                 int64
Product Category    object
Quantity            int64
Price per Unit      int64
Total Amount        int64
dtype: object
```

```
data.nunique()
```



```
Transaction ID      1000
Date                345
Customer ID         1000
Gender              2
Age                 47
Product Category    3
Quantity            4
```

```

Price per Unit      5
Total Amount       18
dtype: int64

```

```
data.info()
```

```

↳ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Transaction ID        1000 non-null   int64
 1   Date                  1000 non-null   object
 2   Customer ID           1000 non-null   object
 3   Gender                1000 non-null   object
 4   Age                   1000 non-null   int64
 5   Product Category      1000 non-null   object
 6   Quantity              1000 non-null   int64
 7   Price per Unit        1000 non-null   int64
 8   Total Amount          1000 non-null   int64
dtypes: int64(5), object(4)
memory usage: 70.4+ KB

```

```
pd.isnull(data).sum()
```

```

↳ Transaction ID      0
Date                  0
Customer ID           0
Gender                0
Age                   0
Product Category      0
Quantity              0
Price per Unit        0
Total Amount          0
dtype: int64

```

```

data[["Date"]] = data[["Date"]].apply(pd.to_datetime)
data["Quantity"] = data["Quantity"].astype(float).astype('Int64')
data.dtypes


```

```

↳ <ipython-input-10-759fd6069785>:1: UserWarning: Parsing dates in %d-%m-%Y format when da
data[["Date"]] = data[["Date"]].apply(pd.to_datetime)
Transaction ID      int64
Date                datetime64[ns]
Customer ID         object
Gender              object
Age                 int64
Product Category    object
Quantity            Int64
Price per Unit      int64
Total Amount        int64
dtype: object


```

```
data.nunique()
```




Transaction ID	1000
Date	345
Customer ID	1000
Gender	2
Age	47
Product Category	3
Quantity	4
Price per Unit	5
Total Amount	18
dtype: int64	

```
data.describe()
```



	Transaction ID	Date	Age	Quantity	Price per Unit	Total Amount
count	1000.000000	1000	1000.000000	1000.0	1000.000000	1000.000000
mean	500.500000	2023-07-03 00:25:55.200000256	41.39200	2.514	179.890000	456.000000
min	1.000000	2023-01-01 00:00:00	18.00000	1.0	25.000000	25.000000
25%	250.750000	2023-04-08 00:00:00	29.00000	1.0	30.000000	60.000000
50%	500.500000	2023-06-29 12:00:00	42.00000	3.0	50.000000	135.000000
75%	750.250000	2023-10-04 00:00:00	53.00000	4.0	300.000000	900.000000
max	1000.000000	2024-01-01 00:00:00	64.00000	4.0	500.000000	2000.000000

```
#DATA_CLEANING
data.isnull().sum()
```



Transaction ID	0
Date	0
Customer ID	0
Gender	0
Age	0
Product Category	0
Quantity	0
Price per Unit	0
Total Amount	0
dtype: int64	

```
#null values are set to 1
data['Quantity'].fillna(1, inplace=True)
```

```
data['Price per Unit'].fillna(data.groupby('Product Category')['Price per Unit'].transform('
data['Price per Unit'].fillna(data['Price per Unit'].mean(), inplace=True)
```

```
data = data[data['Product Category'].notna()]
```

```
data.loc[data['Date'] > data['Date'], 'Date'] = 'Past Due'
```

```
print(data['Total Amount'].mean(), data['Total Amount'].std(), data['Total Amount'].mean() -
data = data[data['Total Amount'] < (data['Total Amount'].mean() + (3 * data['Total Amount'].
data = data[data['Total Amount'] > (data['Total Amount'].mean() - (3 * data['Total Amount'].)
```

```
⇒ 456.0 559.997631555123 -1223.992894665369 2135.992894665369
```


```
data.groupby(['Product Category']).agg({'Quantity':np.sum}).reset_index()
```

```
⇒
```

	Product Category	Quantity
0	Beauty	771
1	Clothing	894
2	Electronics	849

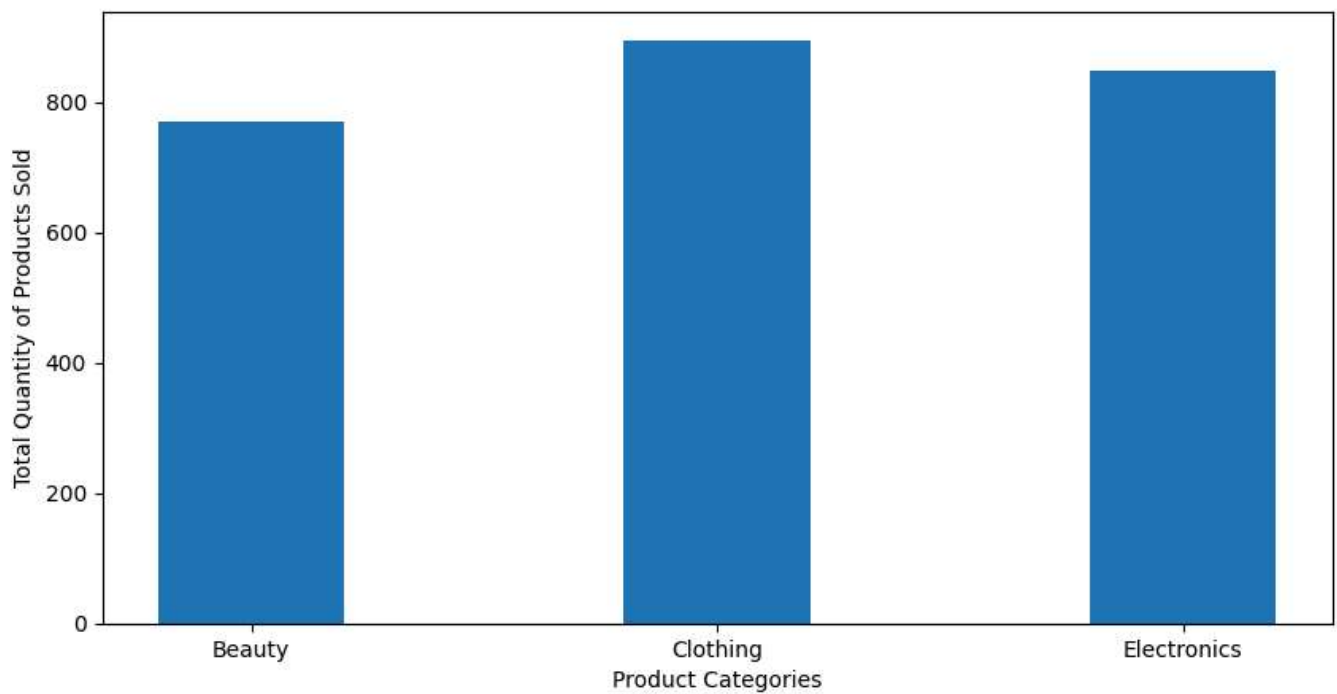
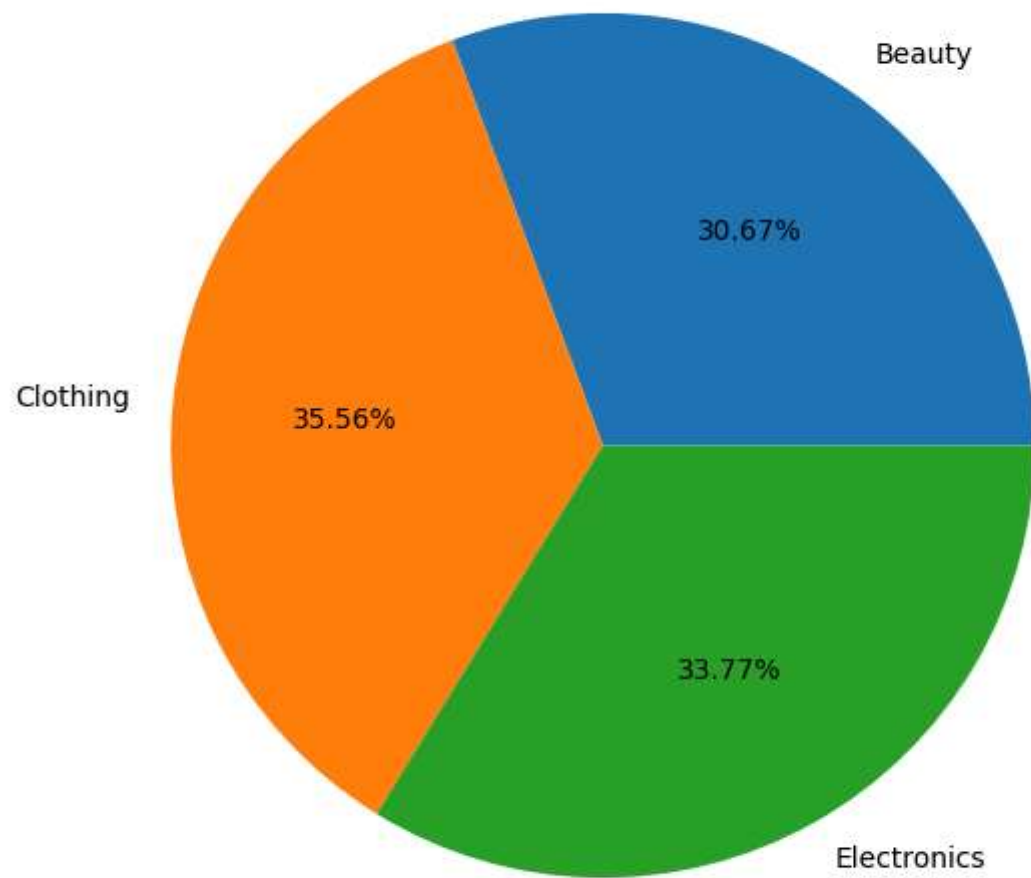
```
plt.figure(figsize =(10, 12))
sns.boxplot(y= data['Price per Unit'], x = data['Product Category'])
```

↔ <Axes: xlabel='Product Category', ylabel='Price per Unit'>







```
df = data.groupby(['Product Category']).agg({'Quantity':np.sum}).reset_index()
fig1 = plt.figure(figsize =(10, 7))
plt.pie(df.Quantity, labels = df['Product Category'],autopct='%1.2f%%')
# show plot
plt.show()


fig2 = plt.figure(figsize = (10, 5))
plt.bar(df['Product Category'], df.Quantity, width = 0.4)
plt.xlabel("Product Categories")
plt.ylabel("Total Quantity of Products Sold")
# show plot
plt.show()
```





```
#Plot histogram of all numeric attrubites to see their distribution
# Plot the histograms of each
data.hist(bins=30, figsize=(30,20))
plt.show()
```



```
# Product Category
df=df.sort_values('Quantity')
fig2 = plt.figure(figsize = (10, 5))
plt.bar(df['Product Category'], df.Quantity, color = 'maroon', width = 0.4)
plt.xlabel("Product Categories")
plt.ylabel("Total Quantity of Products Sold")
# show plot
plt.show()
```



```
#Region
df = data.groupby(['Gender']).agg({'Quantity':np.sum}).reset_index().sort_values('Quantity')
fig = plt.figure(figsize = (10, 5))
plt.bar(df.Gender, df.Quantity, color = 'maroon', width = 0.6)
plt.xlabel("Gender")
plt.ylabel("Total Quantity of Products Sold")
# show plot
plt.show()
```



```
#Plot number of transactions on each day of week.
```

```
df = data.groupby(['Date']).size()
```

```
new_df = df.to_frame(name = 'size').reset_index()
```

```
new_df['NumberofTransactions']=1
```

```
new_df['day_of_week'] = new_df['Date'].dt.day_name()
```

```
dataTransactions = new_df.groupby('day_of_week')['NumberofTransactions'].agg('sum').reset_index()
```

```
dataTransactions = dataTransactions.loc[[0,1,2,3,4,5,6], :] # Sunday to Saturday
```

```
# #plotting bar chart
```


```
fig = plt.figure(figsize = (15, 5))
```

```
plt.bar(dataTransactions.day_of_week, dataTransactions.NumberofTransactions, color = 'maroon')
```

```
plt.xlabel("Day of week")
```

```
plt.ylabel("Total transactions")
```

```
plt.show()
```



```
df.head(20)
```

```
df = data.groupby(['Transaction ID'], sort=False).size().reset_index(name='Count')
```

```
df['Count'].describe()
```



```
count    1000.0
mean      1.0
std       0.0
min       1.0
25%       1.0
50%       1.0
75%       1.0
max       1.0
Name: Count, dtype: float64
```

```
data['NormalizedPrice'] = (data['Price per Unit'] - data['Price per Unit'].mean()) / data['Price per Unit']
data.head(1000)
```



	Transaction ID	Date	Customer ID	Gender	Age	Product Category	Quantity	Price per Unit	Total Amount	Number of Transactions
0	1	2023-11-24 00:00:00	CUST001	Male	34	Beauty	3	50	150	1
1	2	2023-02-27 00:00:00	CUST002	Female	26	Clothing	2	500	1000	1
2	3	2023-01-13 00:00:00	CUST003	Male	50	Electronics	1	30	30	1

```
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import norm
```

```
plt.hist(data['Price per Unit'], bins=20)
plt.xlabel('Price per Unit')
plt.ylabel('Frequency')
plt.title('Histogram of Price per Unit')
```

```
# Calculate the mean and standard deviation of the column
mean = np.mean(data['Price per Unit'])
std = np.std(data['Price per Unit'])
```

```
# Calculate the theoretical normal distribution
xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
p = norm.pdf(x, mean, std)
```

```
# Plot the theoretical normal distribution on top of the histogram
plt.plot(x, p, 'k', linewidth=2)
plt.show()
```

